

# IPv4

From Wikipedia, the free encyclopedia

**Internet Protocol version 4** (**IPv4**) is the fourth version of the [Internet Protocol](#) (IP). It is one of the core protocols of standards-based [internetworking](#) methods in the [Internet](#) and other [packet-switched](#) networks. IPv4 was the first version deployed for production on [SATNET](#) in 1982 and on the [ARPANET](#) in January 1983. It is still used to route most Internet traffic today,<sup>[1]</sup> despite the ongoing deployment of a successor protocol, [IPv6](#).

IPv4 uses a 32-bit address space which provides 4,294,967,296 (2<sup>32</sup>) unique addresses, but large blocks are reserved for special networking methods.

## Contents

- 1

History
- 2

Purpose
- 3

Addressing

3.1

Address representations

3.2

Allocation

3.3

Special-use addresses

3.3.1

Private networks

3.4

Link-local addressing

3.5

Loopback

3.6

First and last subnet addresses

3.7

Address resolution

4

Address space exhaustion

5

Packet structure

5.1

Header

5.2

Data

6

Fragmentation and reassembly

6.1

Fragmentation

6.2

Reassembly

7

Assistive protocols

8

See also

9

Notes

10

References

11

External links

## History

This section **needs expansion**.  
You can help by [adding to it](#).  
*(August 2020)*

The IP layer was originally separated in the v3 of the TCP for design improvement, and stabilised in version 4.<sup>[2]</sup> IPv4 is described in [IETF](#) publication RFC 791 (September 1981), replacing an earlier definition (RFC 760, January 1980). In March 1982, the US Department of Defense declared TCP/IP as the standard for all military computer networking.<sup>[3]</sup>

## Purpose

The Internet Protocol is the protocol that defines and enables [internetworking](#) at the [internet layer](#) of the [Internet Protocol Suite](#). In essence it forms the Internet. It uses a logical addressing system and performs *routing*, which is the forwarding of packets from a source host to the next router that is one hop closer to the intended destination host on another network.

IPv4 is a [connectionless](#) protocol, and operates on a [best-effort delivery](#) model, in that it does not guarantee delivery, nor does it assure proper sequencing or avoidance of duplicate delivery. These aspects, including data integrity, are addressed by an [upper layer](#) transport protocol, such as the [Transmission Control Protocol](#) (TCP).

## Addressing

IPv4 uses 32-bit addresses which limits the [address space](#) to 4 294 967 296 (2<sup>32</sup>) addresses.

IPv4 reserves special address blocks for [private networks](#) (~18 million addresses) and [multicast](#) addresses (~270 million addresses).

### Address representations

IPv4 addresses may be represented in any notation expressing a 32-bit integer value. They are most often written in [dot-decimal notation](#), which consists of four [octets](#) of the address expressed individually in [decimal](#) numbers and separated by [periods](#).

For example, the quad-dotted IP address 192.0.2.235 represents the 32-bit [decimal](#) number 3221226219, which in [hexadecimal](#) format is 0xC00002EB. This may also be expressed in dotted hex format as 0xC0.0x00.0x02.0xEB, or with octal byte values as 0300.0000.0002.0353.

[CIDR notation](#) combines the address with its routing prefix in a compact format, in which the address is followed by a slash character (/) and the count of leading consecutive 1 bits in the routing prefix (subnet mask).

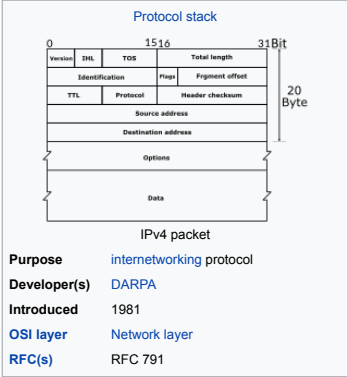
Other address representations were in common use when [classful networking](#) was practiced. For example, the loopback address 127.0.0.1 is commonly written as 127.1, given that it belongs to a class-A network with eight bits for the network mask and 24 bits for the host number. When fewer than four numbers are specified in the address in dotted notation, the last value is treated as an integer of as many bytes as are required to fill out the address to four octets. Thus, the address 127.65530 is equivalent to 127.0.255.250.

### Allocation

In the original design of IPv4, an IP address was divided into two parts: the network identifier was the most significant octet of the address, and the host identifier was the rest of the address. The latter was also called the *rest field*. This structure permitted a maximum of 256 network identifiers, which was quickly found to be inadequate.

To overcome this limit, the most-significant address octet was redefined in 1981 to create *network classes*, in a system which later became known as [classful networking](#). The revised system defined five classes. Classes A, B, and C had different bit lengths for network identification. The rest of the address was used as previously to identify a host within a network.

### Internet Protocol version 4



Because of the different sizes of fields in different classes, each network class had a different capacity for addressing hosts. In addition to the three classes for addressing hosts, Class D was defined for **multicast** addressing and Class E was reserved for future applications.

Dividing existing classful networks into subnets began in 1985 with the publication of [RFC 950](#). This division was made more flexible with the introduction of variable-length subnet masks (VLSM) in [RFC 1109](#) in 1987. In 1993, based on this work, [RFC 1517](#) introduced **Classless Inter-Domain Routing** (CIDR),<sup>[4]</sup> which expressed the number of bits (from the **most significant**) as, for instance, /24, and the class-based scheme was dubbed *classful*, by contrast. CIDR was designed to permit repartitioning of any address space so that smaller or larger blocks of addresses could be allocated to users. The hierarchical structure created by CIDR is managed by the [Internet Assigned Numbers Authority](#) (IANA) and the [regional Internet registries](#) (RIRs). Each RIR maintains a publicly searchable [WHOIS](#) database that provides information about IP address assignments.

**Special-use addresses** [[edit](#)]

The [Internet Engineering Task Force](#) (IETF) and IANA have restricted from general use various [reserved IP addresses](#) for special purposes.<sup>[5]</sup> Notably these addresses are used for **multicast** traffic and to provide addressing space for unrestricted uses on private networks.

Special address blocks				
Address block	Address range	Number of addresses	Scope	Description
0.0.0.0/8	0.0.0.0–0.255.255.255	16 777 216	Software	Current network <sup>[6]</sup> (only valid as source address).
10.0.0.0/8	10.0.0.0–10.255.255.255	16 777 216	Private network	Used for local communications within a <a href="#">private network</a> . <sup>[7]</sup>
100.64.0.0/10	100.64.0.0–100.127.255.255	4 194 304	Private network	<a href="#">Shared address space</a> <sup>[8]</sup> for communications between a service provider and its subscribers when using a <a href="#">carrier-grade NAT</a> .
127.0.0.0/8	127.0.0.0–127.255.255.255	16 777 216	Host	Used for <a href="#">loopback addresses</a> to the local host. <sup>[6]</sup>
169.254.0.0/16	169.254.0.0–169.254.255.255	65 536	Subnet	Used for <a href="#">link-local addresses</a> <sup>[9]</sup> between two hosts on a single link when no IP address is otherwise specified, such as would have normally been retrieved from a <a href="#">DHCP</a> server.
172.16.0.0/12	172.16.0.0–172.31.255.255	1 048 576	Private network	Used for local communications within a private network. <sup>[7]</sup>
192.0.0.0/24	192.0.0.0–192.0.0.255	256	Private network	IETF Protocol Assignments. <sup>[6]</sup>
192.0.2.0/24	192.0.2.0–192.0.2.255	256	Documentation	Assigned as TEST-NET-1, documentation and examples. <sup>[10]</sup>
192.88.99.0/24	192.88.99.0–192.88.99.255	256	Internet	Reserved. <sup>[11]</sup> Formerly used for <a href="#">IPv6 to IPv4 relay</a> . <sup>[12]</sup> (included <a href="#">IPv6</a> address block <a href="#">2002::/16</a> ).
192.168.0.0/16	192.168.0.0–192.168.255.255	65 536	Private network	Used for local communications within a private network. <sup>[7]</sup>
198.18.0.0/15	198.18.0.0–198.19.255.255	131 072	Private network	Used for benchmark testing of inter-network communications between two separate subnets. <sup>[13]</sup>
198.51.100.0/24	198.51.100.0–198.51.100.255	256	Documentation	Assigned as TEST-NET-2, documentation and examples. <sup>[10]</sup>
203.0.113.0/24	203.0.113.0–203.0.113.255	256	Documentation	Assigned as TEST-NET-3, documentation and examples. <sup>[10]</sup>
224.0.0.0/4	224.0.0.0–239.255.255.255	268 435 456	Internet	In use for <a href="#">IP multicast</a> . <sup>[14]</sup> (Former Class D network).
233.252.0.0/24	233.252.0.0–233.252.0.255	256	Documentation	Assigned as MCAST-TEST-NET, documentation and examples. <sup>[14][15]</sup>
240.0.0.0/4	240.0.0.0–255.255.255.254	268 435 455	Internet	Reserved for future use. <sup>[16]</sup> (Former Class E network).
255.255.255.255/32	255.255.255.255	1	Subnet	Reserved for the "limited <a href="#">broadcast</a> " destination address. <sup>[6][17]</sup>

**Private networks** [[edit](#)]

Of the approximately four billion addresses defined in IPv4, about 18 million addresses in three ranges are reserved for use in [private networks](#). Packets addresses in these ranges are not routable in the public Internet; they are ignored by all public routers. Therefore, private hosts cannot directly communicate with public networks, but require [network address translation](#) at a routing gateway for this purpose.

Reserved private IPv4 network ranges <sup>[7]</sup>				
Name	CIDR block	Address range	Number of addresses	<i>Classful</i> description
24-bit block	10.0.0.0/8	10.0.0.0 – 10.255.255.255	16 777 216	Single Class A.
20-bit block	172.16.0.0/12	172.16.0.0 – 172.31.255.255	1 048 576	Contiguous range of 16 Class B blocks.
16-bit block	192.168.0.0/16	192.168.0.0 – 192.168.255.255	65 536	Contiguous range of 256 Class C blocks.

Since two private networks, e.g., two branch offices, cannot directly interoperate via the public Internet, the two networks must be bridged across the Internet via a [virtual private network](#) (VPN) or an [IP tunnel](#), which [encapsulates](#) packets, including their headers containing the private addresses, in a protocol layer during transmission across the public network. Additionally, encapsulated packets may be encrypted for transmission across public networks to secure the data.

**Link-local addressing** [[edit](#)]

RFC 3927 defines the special address block 169.254.0.0/16 for link-local addressing. These addresses are only valid on the link (such as a local network segment or point-to-point connection) directly connected to a host that uses them. These addresses are not routable. Like private addresses, these addresses cannot be the source or destination of packets traversing the internet. These addresses are primarily used for address autoconfiguration ([Zeroconf](#)) when a host cannot obtain an IP address from a DHCP server or other internal configuration methods.

When the address block was reserved, no standards existed for address autoconfiguration. [Microsoft](#) created an implementation called [Automatic Private IP Addressing](#) (APIPA), which was deployed on millions of machines and became a [de facto standard](#). Many years later, in May 2005, the [IETF](#) defined a formal standard in RFC 3927, entitled *Dynamic Configuration of IPv4 Link-Local Addresses*.

**Loopback** [[edit](#)]

*Main article: [Localhost](#)*

The class A network 127.0.0.0 (classless network 127.0.0.0/8) is reserved for **loopback**. IP packets whose source addresses belong to this network should never appear outside a host. Packets received on a non-loopback interface with a loopback source or destination address must be dropped.

**First and last subnet addresses** [ edit ]

See also: *IPv4 subnetting reference*

The first address in a subnet is used to identify the subnet itself. In this address all host bits are *0*. To avoid ambiguity in representation, this address is reserved.<sup>[18]</sup> The last address has all host bits set to *1*. It is used as a local **broadcast address** for sending messages to all devices on the subnet simultaneously. For networks of size /24 or larger, the broadcast address always ends in 255.

For example, in the subnet 192.168.5.0/24 (subnet mask 255.255.255.0) the identifier 192.168.5.0 is used to refer to the entire subnet. The broadcast address of the network is 192.168.5.255.

	Binary form	Dot-decimal notation
Network space	11000000.10101000.00000101.00000000	192.168.5.0
Broadcast address	11000000.10101000.00000101.11111111	192.168.5.255
In red, is shown the host part of the IP address; the other part is the network prefix. The host gets inverted (logical NOT), but the network prefix remains intact.		

However, this does not mean that every address ending in 0 or 255 cannot be used as a host address. For example, in the /16 subnet 192.168.0.0/255.255.0.0, which is equivalent to the address range 192.168.0.0–192.168.255.255, the broadcast address is 192.168.255.255. One can use the following addresses for hosts, even though they end with 255: 192.168.1.255, 192.168.2.255, etc. Also, 192.168.0.0 is the network identifier and must not be assigned to an interface.<sup>[19]</sup> The addresses 192.168.1.0, 192.168.2.0, etc., may be assigned, despite ending with 0.

In the past, conflict between network addresses and broadcast addresses arose because some software used non-standard broadcast addresses with zeros instead of ones.<sup>[20]</sup>

In networks smaller than /24, broadcast addresses do not necessarily end with 255. For example, a CIDR subnet 203.0.113.16/28 has the broadcast address 203.0.113.31.

	Binary form	Dot-decimal notation
Network space	11001011.00000000.01110001.00010000	203.0.113.16
Broadcast address	11001011.00000000.01110001.00011111	203.0.113.31
In red, is shown the host part of the IP address; the other part is the network prefix. The host gets inverted (logical NOT), but the network prefix remains intact.		

As a special case, a /31 network has capacity for just two hosts. These networks are typically used for point-to-point connections. There is no network identifier or broadcast address for these networks.<sup>[21]</sup>

**Address resolution** [ edit ]

Main article: *Domain Name System*

Hosts on the **Internet** are usually known by names, e.g., www.example.com, not primarily by their IP address, which is used for routing and network interface identification. The use of domain names requires translating, called *resolving*, them to addresses and vice versa. This is analogous to looking up a phone number in a phone book using the recipient's name.

The translation between addresses and domain names is performed by the **Domain Name System** (DNS), a hierarchical, distributed naming system that allows for the subdelegation of **namespaces** to other DNS servers.

**Address space exhaustion** [ edit ]

Main article: *IPv4 address exhaustion*

Since the 1980s, it was apparent that the pool of available IPv4 addresses was depleting at a rate that was not initially anticipated in the original design of the network.<sup>[22]</sup> The main market forces that accelerated address depletion included the rapidly growing number of Internet users, who increasingly used mobile computing devices, such as **laptop computers**, **personal digital assistants** (PDAs), and **smart phones** with IP data services. In addition, high-speed Internet access was based on always-on devices. The threat of exhaustion motivated the introduction of a number of remedial technologies, such as **Classless Inter-Domain Routing** (CIDR) methods by the mid-1990s, pervasive use of **network address translation** (NAT) in network access provider systems, and strict usage-based allocation policies at the regional and local Internet registries.

The primary address pool of the Internet, maintained by IANA, was exhausted on 3 February 2011, when the last five blocks were allocated to the five **RIRs**.<sup>[23]</sup><sup>[24]</sup> **APNIC** was the first RIR to exhaust its regional pool on 15 April 2011, except for a small amount of address space reserved for the transition technologies to IPv6, which is to be allocated under a restricted policy.<sup>[25]</sup>

The long-term solution to address exhaustion was the 1998 specification of a new version of the Internet Protocol, **IPv6**.<sup>[26]</sup> It provides a vastly increased address space, but also allows improved route aggregation across the Internet, and offers large subnetwork allocations of a minimum of 2<sup>64</sup> host addresses to end users. However, IPv4 is not directly interoperable with IPv6, so that IPv4-only hosts cannot directly communicate with IPv6-only hosts. With the phase-out of the **6bone** experimental network starting in 2004, permanent formal deployment of IPv6 commenced in 2006.<sup>[27]</sup> Completion of **IPv6 deployment** is expected to take considerable time,<sup>[28]</sup> so that intermediate **transition technologies** are necessary to permit hosts to participate in the Internet using both versions of the protocol.

**Packet structure** [ edit ]

An IP **packet** consists of a header section and a data section. An IP packet has no data checksum or any other footer after the data section. Typically the **link layer** encapsulates IP packets in frames with a CRC footer that detects most errors, many **transport-layer protocols** carried by IP also have their own error checking.<sup>[29]</sup>

**Header** [ edit ]

The IPv4 packet header consists of 14 fields, of which 13 are required. The 14th field is optional and aptly named: options. The fields in the header are packed with the most significant byte first (**big endian**), and for the diagram and discussion, the most significant bits are considered to come first (**MSB 0 bit numbering**). The most significant bit is numbered 0, so the version field is actually found in the four most significant bits of the first byte, for example.

IPv4 header format																																		
Offsets	Octet	0								1								2								3								
Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	
0	0	Version				IHL				DSCP						ECN				Total Length														
4	32	Identification															Flags			Fragment Offset														
8	64	Time To Live								Protocol								Header Checksum																
12	96	Source IP Address																																
16	128	Destination IP Address																																
20	160	Options (if IHL > 5)																																
⋮	⋮																																	
60	480																																	

Version

Version

The first header field in an IP [packet](#) is the four-bit version field. For IPv4, this is always equal to 4.

**Internet Header Length (IHL)**

The IPv4 header is variable in size due to the optional 14th field (options). The IHL field contains the size of the IPv4 header, it has 4 bits that specify the number of 32-bit words in the header. The minimum value for this field is 5,<sup>[30]</sup> which indicates a length of  $5 \times 32 \text{ bits} = 160 \text{ bits} = 20 \text{ bytes}$ . As a 4-bit field, the maximum value is 15, this means that the maximum size of the IPv4 header is  $15 \times 32 \text{ bits} = 480 \text{ bits} = 60 \text{ bytes}$ .

**Differentiated Services Code Point (DSCP)**

Originally defined as the [type of service](#) (ToS), this field specifies [differentiated services](#) (DiffServ) per RFC 2474.<sup>[a]</sup> Real-time data streaming makes use of the DSCP field. An example is [Voice over IP](#) (VoIP), which is used for interactive voice services.

**Explicit Congestion Notification (ECN)**

This field is defined in RFC 3168 and allows end-to-end notification of [network congestion](#) without [dropping packets](#). ECN is an optional feature available when both endpoints support it and effective when also supported by the underlying network.

**Total Length**

This 16-bit field defines the entire packet size in bytes, including header and data. The minimum size is 20 bytes (header without data) and the maximum is 65,535 bytes. All hosts are required to be able to reassemble datagrams of size up to 576 bytes, but most modern hosts handle much larger packets. Links may impose further restrictions on the packet size, in which case datagrams must be [fragmented](#). Fragmentation in IPv4 is performed in either the sending host or in routers. Reassembly is performed at the receiving host.

**Identification**

This field is an identification field and is primarily used for uniquely identifying the group of fragments of a single IP datagram. Some experimental work has suggested using the ID field for other purposes, such as for adding packet-tracing information to help trace datagrams with spoofed source addresses,<sup>[31]</sup> but RFC 6864 now prohibits any such use.

**Flags**

A three-bit field follows and is used to control or identify fragments. They are (in order, from most significant to least significant):

- bit 0: Reserved; must be zero.<sup>[b]</sup>
- bit 1: Don't Fragment (DF)
- bit 2: More Fragments (MF)

If the DF flag is set, and fragmentation is required to route the packet, then the packet is dropped. This can be used when sending packets to a host that does not have resources to perform reassembly of fragments. It can also be used for [path MTU discovery](#), either automatically by the host IP software, or manually using diagnostic tools such as [ping](#) or [traceroute](#).

For unfragmented packets, the MF flag is cleared. For fragmented packets, all fragments except the last have the MF flag set. The last fragment has a non-zero Fragment Offset field, differentiating it from an unfragmented packet.

**Fragment Offset**

This field specifies the offset of a particular fragment relative to the beginning of the original unfragmented IP datagram in units of eight-byte blocks. The first fragment has an offset of zero. The 13 bit field allows a maximum offset of  $(2^{13} - 1) \times 8 = 65,528 \text{ bytes}$ , which, with the header length included ( $65,528 + 20 = 65,548 \text{ bytes}$ ), supports fragmentation of packets exceeding the maximum IP length of 65,535 bytes.

**Time To Live (TTL)**

An eight-bit [time to live](#) field helps prevent datagrams from persisting (e.g. going in circles) on an internet. This field limits a datagram's lifetime. It is specified in seconds, but time intervals less than 1 second are rounded up to 1. In practice, the field has become a [hop count](#)—when the datagram arrives at a [router](#), the router decrements the TTL field by one. When the TTL field hits zero, the router discards the packet and typically sends an [ICMP Time Exceeded](#) message to the sender.

The program [traceroute](#) uses these ICMP Time Exceeded messages to print the routers used by packets to go from the source to the destination.

**Protocol**

This field defines the protocol used in the data portion of the IP datagram. IANA maintains a [list of IP protocol numbers](#) as directed by RFC 790.

**Header Checksum**

The 16-bit [IPv4 header checksum](#) field is used for error-checking of the header. When a packet arrives at a router, the router calculates the checksum of the header and compares it to the checksum field. If the values do not match, the router discards the packet. Errors in the data field must be handled by the encapsulated protocol. Both [UDP](#) and [TCP](#) have checksum fields.

When a packet arrives at a router, the router decreases the TTL field. Consequently, the router must calculate a new checksum.

**Source address**

This field is the [IPv4 address](#) of the sender of the packet. Note that this address may be changed in transit by a [network address translation](#) device.

**Destination address**

This field is the [IPv4 address](#) of the receiver of the packet. As with the source address, this may be changed in transit by a [network address translation](#) device.

**Options**

The [options field](#) is not often used. Note that the value in the IHL field must include enough extra 32-bit words to hold all the options (plus any padding needed to ensure that the header contains an integer number of 32-bit words). The list of options may be terminated with an EOL (End of Options List, 0x00) option; this is only necessary if the end of the options would not otherwise coincide with the end of the header. The possible options that can be put in the header are as follows:

Field	Size (bits)	Description
Copied	1	Set to 1 if the options need to be copied into all fragments of a fragmented packet.
Option Class	2	A general options category. 0 is for "control" options, and 2 is for "debugging and measurement". 1 and 3 are reserved.
Option Number	5	Specifies an option.
Option Length	8	Indicates the size of the entire option (including this field). This field may not exist for simple options.
Option Data	Variable	Option-specific data. This field may not exist for simple options.

- Note: If the header length is greater than 5 (i.e., it is from 6 to 15) it means that the options field is present and must be considered.
- Note: Copied, Option Class, and Option Number are sometimes referred to as a single eight-bit field, the *Option Type*.

Packets containing [some options may be considered as dangerous](#) by some routers and be blocked.<sup>[32]</sup>

The table below shows the defined options for IPv4. Strictly speaking, the column labeled "Option Number" is actually the "Option Value" that is derived from the Copied, Option Class, and Option Number bits as defined above. However, since most people today refer to this combined bit set as the "option number," this table shows that common usage. The table shows both the decimal and the hexadecimal option numbers.<sup>[33]</sup>

Option Number	Option Name	Description
0 / 0x00	EOOL	End of Option List
1 / 0x01	NOP	No Operation
2 / 0x02	SEC	Security (defunct)
7 / 0x07	RR	Record Route
10 / 0x0A	ZSU	Experimental Measurement
11 / 0x0B	MTUP	MTU Probe
12 / 0x0C	MTUR	MTU Reply
15 / 0x0F	ENCODE	ENCODE
25 / 0x19	QS	Quick-Start

30 / 0x1E	EXP	RFC3692-style Experiment
68 / 0x44	TS	Time Stamp
82 / 0x52	TR	Traceroute
94 / 0x5E	EXP	RFC3692-style Experiment
130 / 0x82	SEC	Security (RIPSO)
131 / 0x83	LSR	Loose Source Route
133 / 0x85	E-SEC	Extended Security (RIPSO)
134 / 0x86	CIPSO	Commercial IP Security Option
136 / 0x88	SID	Stream ID
137 / 0x89	SSR	Strict Source Route
142 / 0x8E	VISA	Experimental Access Control
144 / 0x90	IMITD	IMI Traffic Descriptor
145 / 0x91	EIP	Extended Internet Protocol
147 / 0x93	ADDEXT	Address Extension
148 / 0x94	RTRALT	Router Alert
149 / 0x95	SDB	Selective Directed Broadcast
151 / 0x97	DPS	Dynamic Packet State
152 / 0x98	UMP	Upstream Multicast Pkt.
158 / 0x9E	EXP	RFC3692-style Experiment
205 / 0xCD	FINN	Experimental Flow Control
222 / 0xDE	EXP	RFC3692-style Experiment

#### Data [[edit](#)]

The packet payload is not included in the checksum. Its contents are interpreted based on the value of the Protocol header field.

Some of the common payload protocols are:

Protocol Number	Protocol Name	Abbreviation
1	<a href="#">Internet Control Message Protocol</a>	ICMP
2	<a href="#">Internet Group Management Protocol</a>	IGMP
6	<a href="#">Transmission Control Protocol</a>	TCP
17	<a href="#">User Datagram Protocol</a>	UDP
41	<a href="#">IPv6 encapsulation</a>	ENCAP
89	<a href="#">Open Shortest Path First</a>	OSPF
132	<a href="#">Stream Control Transmission Protocol</a>	SCTP

See [List of IP protocol numbers](#) for a complete list.

### Fragmentation and reassembly [[edit](#)]

*Main article: [IP fragmentation](#)*

The Internet Protocol enables traffic between networks. The design accommodates networks of diverse physical nature; it is independent of the underlying transmission technology used in the link layer. Networks with different hardware usually vary not only in transmission speed, but also in the [maximum transmission unit](#) (MTU). When one network wants to transmit datagrams to a network with a smaller MTU, it may [fragment](#) its datagrams. In IPv4, this function was placed at the [Internet Layer](#), and is performed in IPv4 routers, which thus require no implementation of any higher layers for the function of routing IP packets.

In contrast, [IPv6](#), the next generation of the Internet Protocol, does not allow routers to perform fragmentation; hosts must determine the path MTU before sending datagrams.

#### Fragmentation [[edit](#)]

When a router receives a packet, it examines the destination address and determines the outgoing interface to use and that interface's MTU. If the packet size is bigger than the MTU, and the Do not Fragment (DF) bit in the packet's header is set to 0, then the router may fragment the packet.

The router divides the packet into fragments. The max size of each fragment is the MTU minus the IP header size (20 bytes minimum; 60 bytes maximum). The router puts each fragment into its own packet, each fragment packet having following changes:

- The *total length* field is the fragment size.
- The *more fragments* (MF) flag is set for all fragments except the last one, which is set to 0.
- The *fragment offset* field is set, based on the offset of the fragment in the original data payload. This is measured in units of 8-byte blocks.
- The *header checksum* field is recomputed.

For example, for an MTU of 1,500 bytes and a header size of 20 bytes, the fragment offsets would be multiples of  $\frac{1,500 - 20}{8} = \mathbf{185}$ . These multiples are 0, 185, 370, 555, 740 etc

It is possible that a packet is fragmented at one router, and that the fragments are further fragmented at another router. For example, a packet of 4,520 bytes, including the 20 bytes of the IP header (without options) is fragmented to two packets on a link with an MTU of 2,500 bytes:

Fragment	Size (bytes)	Header size (bytes)	Data size (bytes)	Flag <i>More fragments</i>	Fragment offset (8-byte blocks)
1	2,500	20	2,480	1	0
2	2,040	20	2,020	0	310


The total data size is preserved: 2,480 bytes + 2,020 bytes = 4,500 bytes. The offsets are **0** and  $\frac{0 + 2,480}{8} = \mathbf{310}$ .


On a link with an MTU of 1,500 bytes, each fragment results in two fragments:

Fragment	Size (bytes)	Header size (bytes)	Data size (bytes)	Flag <i>More fragments</i>	Fragment offset (8-byte blocks)
1	1,500	20	1,480	1	0
2	1,020	20	1,000	1	185



- [Dynamically generated graphs of IPv4 address consumption with predictions of exhaustion dates—Geoff Huston](#)
- [IP addressing in China and the myth of address shortage](#)
- [Countdown of remaining IPv4 available addresses](#) *(estimated)*

 [Media from Wikimedia Commons](#)

 [Resources from Wikiversity](#)

**Authority control**  GND: [4588596-5](#) · MA: [2779233093](#)

Categories: [IPv4](#) | [Internet Standards](#) | [Internet layer protocols](#) | [Network layer protocols](#)